

Sunnie S. Y. Kim

sunniesuhyoung@gmail.com
<https://sunniesuhyoung.github.io>

EMPLOYMENT

- 2026–Now **Microsoft Research**
Senior researcher in the Sociotechnical Systems group (FATE, STAC, SMC)
- 2025–2026 **Apple**
Senior research scientist in the Human-Centered Machine Learning & Responsible AI group
- 2023 **Microsoft Research**
Research intern in the Fairness, Accountability, Transparency & Ethics in AI (FATE) group

EDUCATION

- 2020–2025 **Princeton University**
PhD in Computer Science
Dissertation: *Advancing Responsible AI with Human-Centered Evaluation*
Committee: Olga Russakovsky (adviser), Andrés Monroy-Hernández
Jennifer Wortman Vaughan, Q. Vera Liao, Parastoo Abtahi
- 2019–2020 **Toyota Technological Institute at Chicago**
Visiting student advised by Greg Shakhnarovich
- 2014–2018 **Yale University**
Bachelor of Science in Statistics and Data Science
GPA 3.91/4.00, *magna cum laude*, Distinction in the major
Recognition for outstanding dedication to the department
Senior thesis advised by John Lafferty

HONORS, AWARDS & FELLOWSHIPS

- 2025 CHI 2025 Honorable Mention Award 🏆
- 2025 CHI 2025 Special Recognition for Outstanding Review (2 for Papers, 1 for LBW)
- 2025 FAccT 2025 Doctoral Consortium
- 2024 MIT Rising Stars in EECS Recognition ★
- 2024 Siebel Scholars Award (\$35,000) ★
- 2024 Georgia Tech Doctoral Consortium on Responsible Computing, AI, and Society
- 2024 CHI 2024 Doctoral Consortium
- 2024 Princeton SEAS Travel Grant Award
- 2023 CHI 2023 Honorable Mention Award 🏆
- 2023 SIGCHI Gary Marsden Travel Award
- 2022–2025 NSF Graduate Research Fellowship (\$138,000) ★
- 2022–2023 ML Reproducibility Challenge Outstanding Reviewer Award (×2)
- 2018 Yale Adrian Van Sinderen Book Collecting First Prize (\$1,000)
- 2016 Yale Summer Research Fellowship
- 2014–2018 Korea Presidential Science Scholarship (\$200,000) ★

PAPERS

Preprints

Examining Human-Like Behaviors in LLMs: A Multi-Dimensional Analysis of Model Behaviors, User Factors, and System Prompts

Sunnie S. Y. Kim, Margit Bowler, Leon A Gatys

Measuring and Mitigating Overreliance is Necessary for Building Human-Compatible AI

Lujain Ibrahim, Katherine M. Collins, Sunnie S. Y. Kim, Anka Reuel, Max Lamparth, Kevin Feng, Lama Ahmad, Prajna Soni, Alia El Kattan, Merlin Stein, Siddharth Swaroop, Ilia Sucholutsky, Andrew Strait, Q. Vera Liao, Umang Bhatt

PersonaTeaming: Supporting Persona-Driven Red-Teaming for Generative AI

Wesley Hanwen Deng, Mingxi Yan, Sunnie S. Y. Kim, Akshita Jha, Lauren Wilcox, Ken Holstein, Motahhare Eslami, Leon A Gatys

AI Adoption Across Mission-Driven Organizations

Dalia Ali, Muneeb Ahmed, Hailan Wang, Arfa Khan*, Naira Paola Arnez Jordan*, Sunnie S. Y. Kim, Meet Dilip Muchhala, Anne Kathrin Merkle, Orestis Papakyriakopoulos

Conference and Journal Publications (Peer-Reviewed)

2026

Understanding Annotator Safety Policy with Interpretability

Alex Oesterling, Donghao Ren, Yannick Assogba, Dominik Moritz, Sunnie S. Y. Kim, Leon A Gatys*, Fred Hohman*

ACM Conference on Fairness, Accountability, and Transparency (FAccT)

Presenting Large Language Models as Companions Affects What Mental Capacities People Attribute to Them

Allison Chen, Sunnie S. Y. Kim, Angel Nathaniel Franyutti-Cintron, Amaya Dharmasiri, Kushin Mukherjee, Olga Russakovsky, Judith E. Fan

ACM Conference on Human Factors in Computing Systems (CHI)

2025

Fostering Appropriate Reliance on Large Language Models: The Role of Explanations, Sources, and Inconsistencies

Sunnie S. Y. Kim, Jennifer Wortman Vaughan, Q. Vera Liao, Tania Lombrozo, Olga Russakovsky

ACM Conference on Human Factors in Computing Systems (CHI) 🏆 **Honorable Mention Award**

(Featured in Microsoft's New Future of Work Report and presented at 10+ places through invited and contributed talks)

2024

"I'm Not Sure, But...": Examining the Impact of Large Language Models' Uncertainty Expression on User Reliance and Trust

Sunnie S. Y. Kim, Q. Vera Liao, Mihaela Vorvoreanu, Stephanie Ballard, Jennifer Wortman Vaughan

ACM Conference on Fairness, Accountability, and Transparency (FAccT)

(Featured in Axios, New Scientist, ACM showcase, Microsoft's New Future of Work Report, and the Human-Centered AI Medium publication as "Good Reads in Human-Centered AI")

2023

"Help Me Help the AI": Understanding How Explainability Can Support Human-AI Interaction

Sunnie S. Y. Kim, Elizabeth Anne Watkins, Olga Russakovsky, Ruth Fong, Andrés Monroy-Hernández

ACM Conference on Human Factors in Computing Systems (CHI) 🏆 **Honorable Mention Award**

(One of the top 10 cited CHI papers in 2023–2024 (as of Dec 2024); Featured in the Human-Centered AI Medium publication as "CHI 2023 Editors' Choice"; Invited for talks at multiple AI and HCI conference workshops)

Humans, AI, and Context: Understanding End-Users' Trust in a Real-World Computer Vision Application

Sunnie S. Y. Kim, Elizabeth Anne Watkins, Olga Russakovsky, Ruth Fong, Andrés Monroy-Hernández

ACM Conference on Fairness, Accountability, and Transparency (FAccT)

- Overlooked Factors in Concept-based Explanations: Dataset Choice, Concept Learnability, and Human Capability**
 Vikram V. Ramaswamy, Sunnie S. Y. Kim, Ruth Fong, Olga Russakovsky
IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)
- 2022 **HIVE: Evaluating the Human Interpretability of Visual Explanations**
Sunnie S. Y. Kim, Nicole Meister, Vikram V. Ramaswamy, Ruth Fong, Olga Russakovsky
European Conference on Computer Vision (ECCV)
 (Selected as spotlight and invited for talks at multiple AI and HCI conference workshops)
- Shallow Neural Networks Trained to Detect Collisions Recover Features of Visual Loom-Selective Neurons**
 Baohua Zhou, Zifan Li, Sunnie S. Y. Kim, John Lafferty, Damon A. Clark
eLife (Journal for the biomedical and life sciences)
- 2021 **[Re] Don't Judge an Object by Its Context: Learning to Overcome Contextual Bias**
Sunnie S. Y. Kim, Sharon Zhang, Nicole Meister, Olga Russakovsky
ReScience C (Journal for reproducible replications in computational science)
- Fair Attribute Classification through Latent Space De-biasing**
 Vikram V. Ramaswamy, Sunnie S. Y. Kim, Olga Russakovsky
IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)
 (Featured in Coursera's GANs Specialization course and the MIT Press Book *Foundations of Computer Vision*;
 Invited for talks at multiple AI conference workshops)
- Information-Theoretic Segmentation by Inpainting Error Maximization**
 Pedro Savarese, Sunnie S. Y. Kim, Michael Maire, Gregory Shakhnarovich, David McAllester
IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)
- 2020 **Deformable Style Transfer**
Sunnie S. Y. Kim, Nicholas Kolkin, Jason Salavon, Gregory Shakhnarovich
European Conference on Computer Vision (ECCV)
- 2019 **Which Grades Are Better, A's and C's, or all B's? Effects of Variability in Grades on Mock College Admission Decisions**
 Woo-kyoung Ahn, Sunnie S. Y. Kim, Kristen Kim, Peter K. McNally
Judgment and Decision Making (Journal for the psychology of human judgment and decision making)

Workshop Papers and Extended Abstracts (Lightly Peer-Reviewed)

* indicates equal contribution

- 2025 **PersonaTeaming: Exploring How Introducing Personas Can Improve Automated AI Red-Teaming**
 Wesley Hanwen Deng, Sunnie S. Y. Kim, Akshita Jha, Ken Holstein, Motahhare Eslami, Lauren Wilcox, Leon A Gatys
NeurIPS Workshops on Regulatable ML & LLM Evaluation
- Portraying Large Language Models as Machines, Tools, or Companions Affects What Mental Capacities Humans Attribute to Them**
 Allison Chen, Sunnie S. Y. Kim, Amaya Dharmasiri, Olga Russakovsky, Judith E. Fan
CogSci Proceedings & CHI Extended Abstracts (Late Breaking Work)
- Interactivity x Explainability: Toward Understanding How Interactivity Can Improve Computer Vision Explanations**
 Indu Panigrahi, Sunnie S. Y. Kim*, Amna Liaquat*, Rohan Jinturkar, Olga Russakovsky, Ruth Fong, Parastoo Abtahi
CHI Extended Abstracts (Late Breaking Work)

- 2024 **Establishing Appropriate Trust in AI through Transparency and Explainability**
Sunnie S. Y. Kim
CHI Extended Abstracts (Doctoral Consortium)
- Human-Centered Explainable AI (HCXAI): Reloading Explainability in the Era of Large Language Models (LLMs)**
 Upol Ehsan, Elizabeth Anne Watkins, Philipp Wintersberger, Carina Manger, Sunnie S. Y. Kim, Niels van Berkel, Andreas Riener, Mark O. Riedl
CHI Extended Abstracts (Workshop Proposal)
- Allowing Humans to Interactively Guide Machines Where to Look Does Not Always Improve Human-AI Team's Classification Accuracy**
 Giang Nguyen, Mohammad Reza Taesiri, Sunnie S. Y. Kim, Anh Nguyen
CVPR Workshop on Explainable AI for Computer Vision
- 2023 **Explainable AI for End-Users**
Sunnie S. Y. Kim, Elizabeth Anne Watkins, Olga Russakovsky, Ruth Fong, Andrés Monroy-Hernández
CHI Workshop on Human-Centered Explainable AI
- 2022 **Closing the Creator-Consumer Gap in XAI: A Call for Participatory XAI Design with End-users**
Sunnie S. Y. Kim, Elizabeth Anne Watkins, Olga Russakovsky, Ruth Fong, Andrés Monroy-Hernández
NeurIPS Workshop on Human-Centered AI
- ELUDE: Generating Interpretable Explanations via a Decomposition into Labelled and Unlabelled Features**
 Vikram V. Ramaswamy, Sunnie S. Y. Kim, Nicole Meister, Ruth Fong, Olga Russakovsky
CVPR Workshop on Explainable AI for Computer Vision
- 2021 **Cleaning and Structuring the Label Space of the iMet Collection 2020**
 Vivien Nguyen*, Sunnie S. Y. Kim*
CVPR Workshop on Fine-Grained Visual Categorization

White Papers and Technical Reports (Not Peer-Reviewed)

- 2025 **AI Adoption Across Mission-Driven Organizations**
 Dalia Ali, Muneeb Ahmed, Arfa Khan, Hailan Wang, Sunnie S. Y. Kim, Meet Muchhala, Anne Merkle, Orestis Papakyriakopoulos
TUM Think Tank
- 2018 **Environmental Performance Index**
 Zachary A. Wendling, John W. Emerson, Daniel Esty, Marc Levy, Alex de Sherbinin, ..., Sunnie S. Y. Kim, et al.
World Economic Forum (Environmental Performance Index is a large-scale evaluation of 180 countries' environmental health and ecosystem vitality. As the data team lead, I built the full data pipeline and led the analysis work. The results were presented at the World Economic Forum and covered by international media outlets.)

TALKS

- 2026 Cornell Tech INFO 5390 Fairness in Machine Learning
 Microsoft Research East Coast Labs
- 2025 Princeton COS 436 Human-Computer Interaction
 NSF AI-SDM Workshop on Human-AI Complementarity for Decision Making

	CMU Fairness, Ethics, Accountability, and Transparency Group
	NAVER AI Lab & HCI Group
	Yonsei CSI 7110 Topics in Responsible AI Course
	Princeton COS 598B Machine Behavior Course
	Apple Human-Centered Machine Learning & Responsible AI Group
	Cornell Information Science Colloquium
	Johns Hopkins Computer Science Seminar
	Boston University Computing & Data Sciences Colloquium
	SNU AI Computing Winter School
2024	Cornell Tech Social Technologies Lab
	ECCV 2024 Workshop on Explainable Computer Vision
	Princeton Concepts & Cognition Lab
	MILA Human-Centered AI Reading Group
	IBS Data Science Group
	KAIST Kim Jaechul Graduate School of AI
	NYC Computer Vision Day
2023	Explainable AI Talk Series
	CHI 2023 Workshop on Human-Centered Explainable AI
2022	NeurIPS 2022 Workshop on Human-Centered AI
	CVPR 2022 Workshop on Explainable AI for Computer Vision
2021	CVPR 2021 Workshop on Responsible Computer Vision
	CVPR 2021 Workshop for Women in Computer Vision
2020	Princeton Course COS 429 Computer Vision Course
	Princeton PIXL Talk Series
	Princeton Bias in AI Reading Group

ORGANIZING COMMITTEE

IUI 2027 (Publicity/Web Co-Chair)
FAccT 2025 (Proceedings Co-Chair)
CVPR 2025 Workshop on Explainable AI for Computer Vision (Co-Organizer)
NYC Computer Vision Day 2025 (Event Program Committee)
CVPR 2024 Workshop on Explainable AI for Computer Vision (Co-Organizer)
CHI 2024 Workshop on Human-Centered Explainable AI (Co-Organizer)
CVPR 2023 Workshop on Explainable AI for Computer Vision (Co-Organizer)
CVPR 2023 Workshop for Women in Computer Vision (Co-Organizer)
NESS NextGen Data Science Day 2018 (Local Organizing Committee)

PROGRAM COMMITTEE & REVIEWING

* indicates special recognitions for outstanding reviews

Conferences (Area or Associate Chair)

FAccT 2026 Area Chair

CHI 2026 Associate Chair of the Computational Interaction subcommittee

Conferences (Reviewer)

CHI (2023, 2024, 2025**), FAccT (2023, 2024, 2025), AIES (2024), SaTML (2023)

CVPR (2022, 2023, 2024, 2025), ICCV (2021, 2023), ECCV (2022, 2024)

NeurIPS (2025 Main track & Ethics review, 2026 Main track)

Workshops & Extended Abstracts

IUI 2026 Workshop on Communicating Uncertainty to Foster Realistic Expectations via Human-Centered Design

CVPR 2025 Workshop on Explainable AI for Computer Vision

CHI 2025 Late Breaking Work*

CHI 2024 Workshop on Human-Centered Explainable AI

CVPR 2024 Workshop on Explainable AI for Computer Vision

NeurIPS 2023 Workshop on Explainable AI in Action

ICML 2023 Workshop on AI & HCI

CVPR 2023 Workshop on Explainable AI for Computer Vision

CVPR 2023 Workshop for Women in Computer Vision

AAAI 2023 Workshop on Representation Learning for Responsible Human-Centric AI

CVPR 2021 Workshop on Responsible Computer Vision

Challenges

ML Reproducibility Challenge (2020, 2021*, 2022*)

Books

Foundations of Computer Vision (Authors: Antonio Torralba, Phillip Isola, and William T. Freeman)

Handbook of Human-Centered Artificial Intelligence (Editor-in-Chief: Wei Xu)

MENTORING

Research Mentoring

2025–2026 **Alex Oesterling** (CS PhD student at Harvard)

Understanding Annotator Safety Policy with Interpretability (paper accepted to **FAccT**)

2024–2025 **Allison Chen** (CS PhD student at Princeton. Recipient of the NSF Graduate Research Fellowship)

Understanding How Messages About LLMs Shape People's Mental Capacity Attribution

(papers published in **CHI EA**, **CogSci**, and **CHI**)

2024–2025 **Indu Panigrahi** (CS Master's student at Princeton, now CS PhD student at UIUC)

Incorporating Interactivity in AI Explanations (paper published in **CHI EA**)

- 2022–2023 **Rohan Jinturkar** (CS undergrad at Princeton. Recipient of the Sigma Xi Book Award for Outstanding Undergraduate Research & Outstanding CS Senior Thesis Prize)
Developing an Interactive, Dialogue-based AI Explanation System for Non-Experts (senior thesis)
- 2020–2022 **Nicole Meister** (ECE undergrad at Princeton, now EE PhD student at Stanford. Recipient of the NSF Graduate Research Fellowship, Calvin Dodd MacCracken Senior Thesis/Project Award & Sigma Xi Book Award for Outstanding Undergraduate Research)
Evaluating AI Explanations & Mitigating Contextual Bias in Visual Recognition Systems (papers published in *ECCV* and *ReScience C*)
- 2020–2021 **Sharon Zhang** (Math undergrad at Princeton, now CS PhD student at Stanford. Recipient of the Sigma Xi Book Award for Outstanding Undergraduate Research)
Mitigating Contextual Bias in Visual Recognition Systems (paper published in *ReScience C*)

Non-Research Mentoring

- 2022–2023 Princeton Computer Science G1 Mentoring Program
- 2021–2022 Princeton Computer Science Graduate Applicant Support Program

TEACHING

- 2021 **Princeton Computer Science 429 Computer Vision**
 Graduate Teaching Assistant
- Princeton AI4ALL**
 Instructor
- 2019–2020 **TTI-Chicago Girls Who Code**
 Co-Founder and Instructor
- 2018 **Yale Statistics and Data Science 365/565 Data Mining and Machine Learning**
 Undergraduate Teaching Assistant
- 2017 **Yale Statistics and Data Science 230/530 Data Exploration and Analysis**
 Undergraduate Teaching Assistant

OTHER ACTIVITIES

Community Building

- 2022–2023 Explainable AI Slack and Twitter Community (Co-Organizer)
- 2017–2019 Yale Dimensions Organization for Women and Other Minorities in Math (Co-Founder)

Volunteering

- ECCV (2024), FAccT (2024), CVPR (2022), ICML (2020), ICLR (2020), NeurIPS (2019–2020)
- NSF Safety and Trust in AI-Enabled Systems Workshop (2022)
- COVID Translate Project (2020)

Committee

- 2021 Princeton Computer Science Graduate Admissions Committee
- 2017–2019 Yale Statistics & Data Science Departmental Student Advisory Committee